

# Why you care about relational algebra (even though you didn't know it)

Julian Hyde

@julianhyde

Enterprise Data World

Washington, DC

April 2<sup>nd</sup>, 2015



# About me



Apache  
Calcite



# Why you should care about relational algebra

## Why should you care?

- It is old
- It is as useful as ever
- Exposed in new products such as Hadoop
- New challenges

## Agenda

- Is Hadoop a revolution for the database world?
- What is relational algebra?
- Examples of algebra in action
- Introducing Apache Calcite
- Adding data independence to Hadoop via materialized views

# Old world, new world

## RDBMS



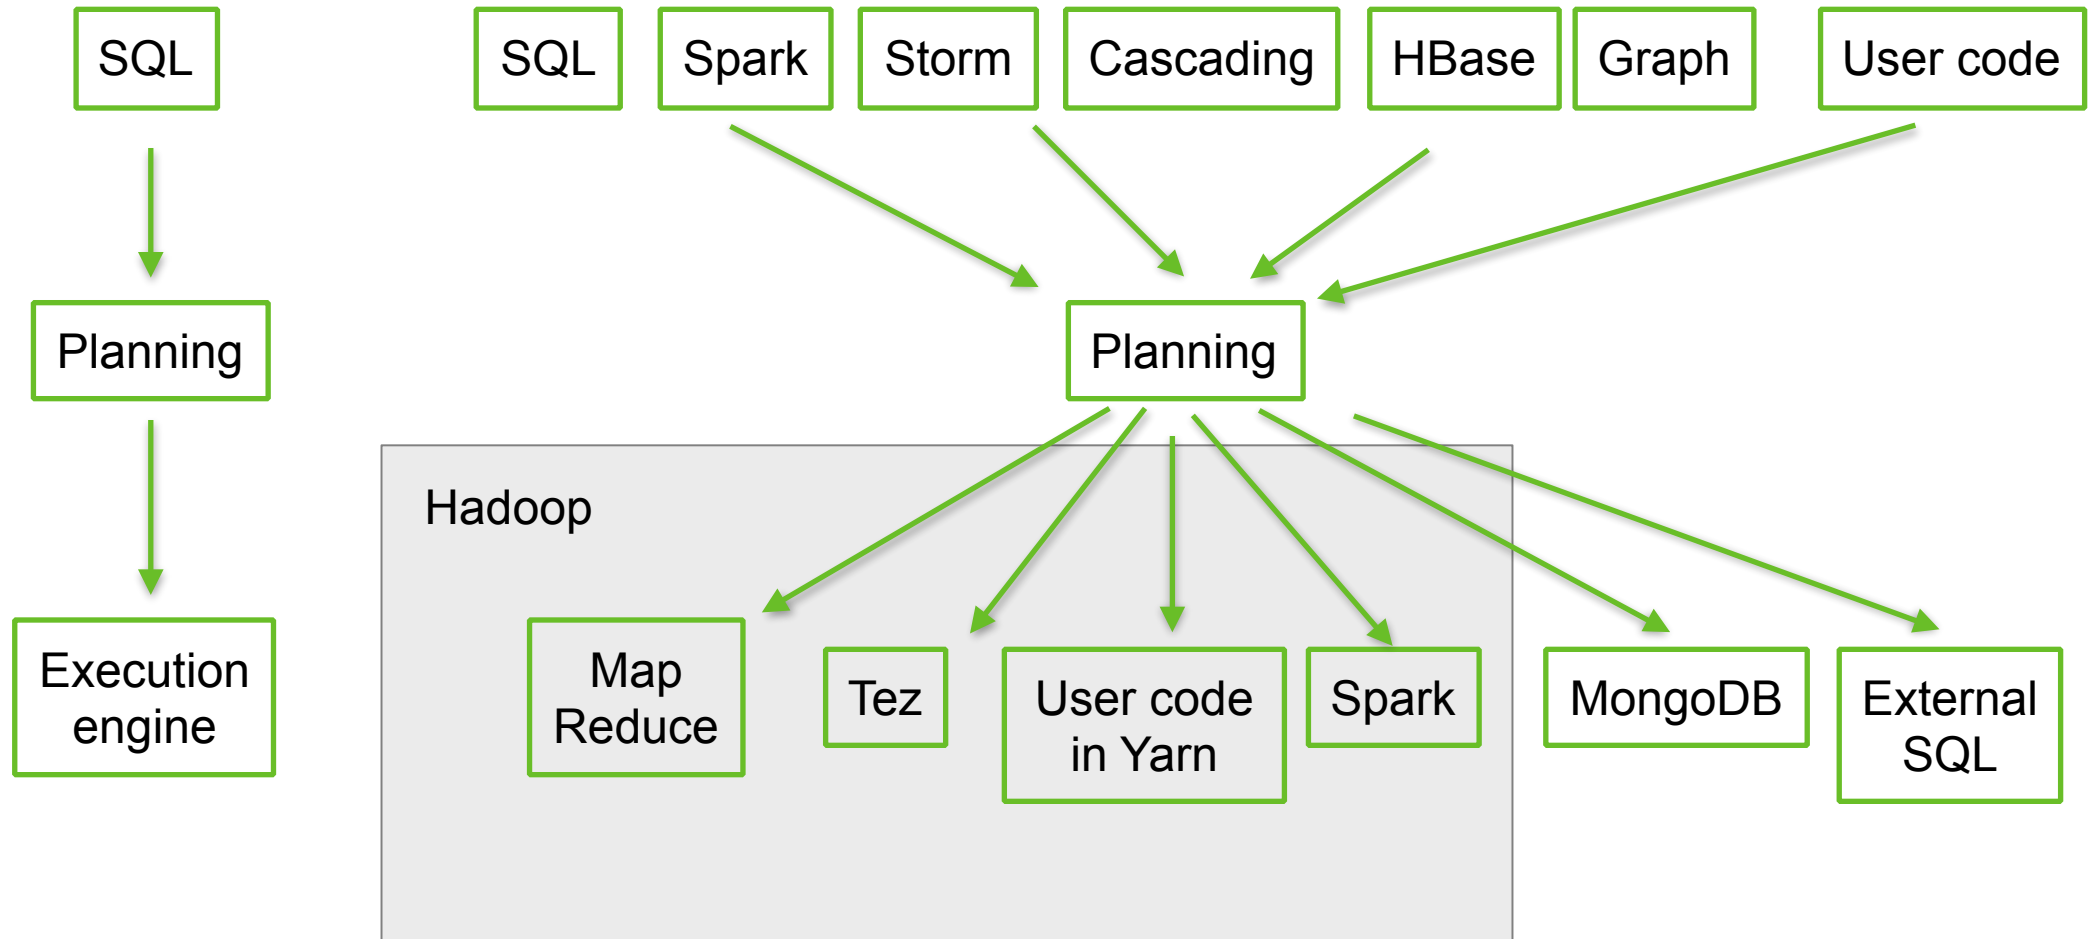
- Security
- Metadata
- SQL
- Query planning
- Data independence

## Hadoop



- Scale
- Late schema
- Choice of front-end
- Choice of engines
- Workload: batch, interactive, streaming, ML, graph, ...


# Many front ends, many engines



# Analogy: LLVM


## Lessons from the compiler community:

- Writing a front end is hard
- Writing a back end is hard
- Writing an optimizer is *really* hard
- Most of the logic in the optimizer is independent of front end and back end
  - E.g. register assignment
- The optimizer is a collection of separate algorithms
- Common language between algorithms



# Relational algebra

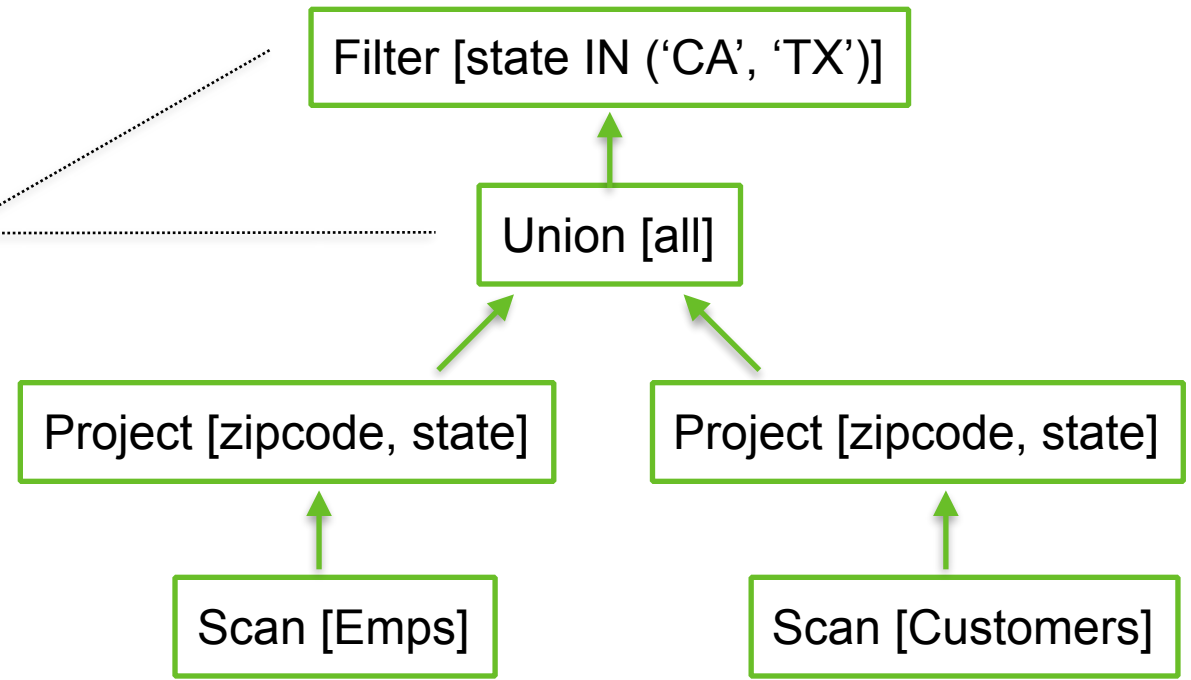
```
SELECT d.name, COUNT(*) AS c
FROM Emps AS e
  JOIN Depts AS d ON e.deptno = d.deptno
WHERE e.age < 30
GROUP BY d.deptno
HAVING COUNT(*) > 5
ORDER BY c DESC
```



(Column names are simplified. They would usually be ordinals, e.g. \$0 is the first column of the left input.)

# Relational algebra - Union and sub-query

```
SELECT * FROM (  
  SELECT zipcode, state  
  FROM Emps  
  UNION ALL  
  SELECT zipcode, state  
  FROM Customers)  
WHERE state IN ('CA', 'TX')
```





# Relational algebra - Insert and Values

**INSERT INTO Facts**  
**VALUES ('Meaning of life', 42),**  
**('Clever as clever', 6)**

Insert [Facts]

Values [['Meaning of life', 42],  
['Clever as clever', 6]]

# Relational algebra - Strict versus Pragmatic

## “Strict” relational algebra

Introduced by E.F. Codd in “A relational model for large shared data banks” [1970]

Goal is mathematical elegance (ability to prove theorems)

Greek symbols:  $\sigma$ ,  $\pi$ ,  $\rho$ ,  $\cup$ ,  $\bowtie$

Relations cannot contain duplicates

Relations are not sorted

Column values are scalars

Only logical operators

## Pragmatic relational algebra

Goal is to optimize queries, allow real-world data models, extensibility

Elegance still important

Verbs: Project, Filter, Union, Join

Relations may contain duplicates

Relations may be sorted

- But Sort is the only logical operator that guarantees order

Null values have 3-value semantics, as in SQL

Physical operators (e.g. HashJoin, MergeJoin)

Physical properties (sort, distribution)

# Algebraic transformations

$(R \text{ filter } c1) \text{ filter } c2 \rightarrow R \text{ filter } (c1 \text{ and } c2)$

$(R1 \text{ union } R2) \text{ join } R3 \text{ on } c \rightarrow (R1 \text{ join } R3 \text{ on } C) \text{ union } (R2 \text{ join } R3 \text{ on } c)$

- Compare distributive law of arithmetic:  $(x + y) * z \rightarrow (x * z) + (y * z)$

$(R1 \text{ join } R2 \text{ on } c) \text{ filter } c2 \rightarrow (R1 \text{ filter } c2) \text{ join } R2 \text{ on } c$  (provided C2 only depends on columns in E, and join is inner)

$(R1 \text{ join } R2 \text{ on } c) \rightarrow (R2 \text{ join } R2 \text{ on } c) \text{ project } [R1.*, R2.*]$

$(R1 \text{ join } R2 \text{ on } c) \text{ join } R3 \text{ on } c2 \rightarrow R1 \text{ join } (R2 \text{ join } R3 \text{ on } c2) \text{ on } c$  (provided c, c2 have the necessary columns)

Many, many others...

# Query using a view

```
SELECT deptno, min(salary)
FROM Managers
WHERE age >= 50
GROUP BY deptno
```

Aggregate [deptno, min(salary)]

Filter [age >= 50]

Scan [Managers]

```
CREATE VIEW Managers AS
SELECT *
FROM Emps
WHERE EXISTS (
  SELECT *
  FROM Emps AS underling
  WHERE underling.manager = emp.id)
```

Project [\$0, \$1, \$2, \$3]

Join [\$0, \$5]

Scan [Emps]


Aggregate [manager]

Scan [Emps]

# After view expansion

```
SELECT deptno, min(salary)
FROM Managers
WHERE age >= 50
GROUP BY deptno
```


```
CREATE VIEW Managers AS
SELECT *
FROM Emps
WHERE EXISTS (
  SELECT *
  FROM Emps AS underling
  WHERE underling.manager = emp.id)
```



# After pushing down filter

```
SELECT deptno, min(salary)
FROM Managers
WHERE age >= 50
GROUP BY deptno
```

```
CREATE VIEW Managers AS
SELECT *
FROM Emps
WHERE EXISTS (
  SELECT *
  FROM Emps AS underling
  WHERE underling.manager = emp.id)
```



# Materialized view

```
CREATE MATERIALIZED VIEW EmpSummary AS  
SELECT deptno,  
       gender,  
       COUNT(*) AS c,  
       SUM(sal) AS s  
FROM Emps  
GROUP BY deptno, gender
```

Scan [EmpSummary]

=

Aggregate [deptno, gender,  
COUNT(\*), SUM(sal)]

Scan [Emps]

Aggregate [COUNT(\*)]

Filter [deptno = 10 AND gender = 'M']

Scan [Emps]

```
SELECT COUNT(*)  
FROM Emps  
WHERE deptno = 10  
AND gender = 'M'
```

# Materialized view, step 2: Rewrite query to match

**CREATE MATERIALIZED VIEW EmpSummary AS**

**SELECT deptno,**

**gender,**

**COUNT(\*) AS c,**

**SUM(sal) AS s**

**FROM Emps**

**GROUP BY deptno, gender**

Scan [EmpSummary]

=

Aggregate [deptno, gender,  
COUNT(\*), SUM(sal)]

Scan [Emps]

Project [c]

Filter [deptno = 10 AND gender = 'M']

**SELECT COUNT(\*)**

**FROM Emps**

**WHERE deptno = 10**

**AND gender = 'M'**

Aggregate [deptno, gender,  
COUNT(\*) AS c, SUM(sal) AS s]

Scan [Emps]



# Materialized view, step 3: Substitute table

**CREATE MATERIALIZED VIEW EmpSummary AS**

**SELECT deptno,**

**gender,**

**COUNT(\*) AS c,**

**SUM(sal) AS s**

**FROM Emps**

**GROUP BY deptno, gender**

Scan [EmpSummary]

=

Aggregate [deptno, gender,  
COUNT(\*), SUM(sal)]

Scan [Emps]

Project [c]

Filter [deptno = 10 AND gender = 'M']

Scan [EmpSummary]

**SELECT COUNT(\*)**

**FROM Emps**

**WHERE deptno = 10**

**AND gender = 'M'**

# Streaming

```
SELECT STREAM DISTINCT productName,  
       floor(rowtime TO HOUR) AS h  
FROM Orders
```

## Delta


Converts a table to a stream

Each time a row is inserted into the table, a record appears in the stream

## Chi

Converts a stream into a table

Often we can safely narrow the table down to a small time window

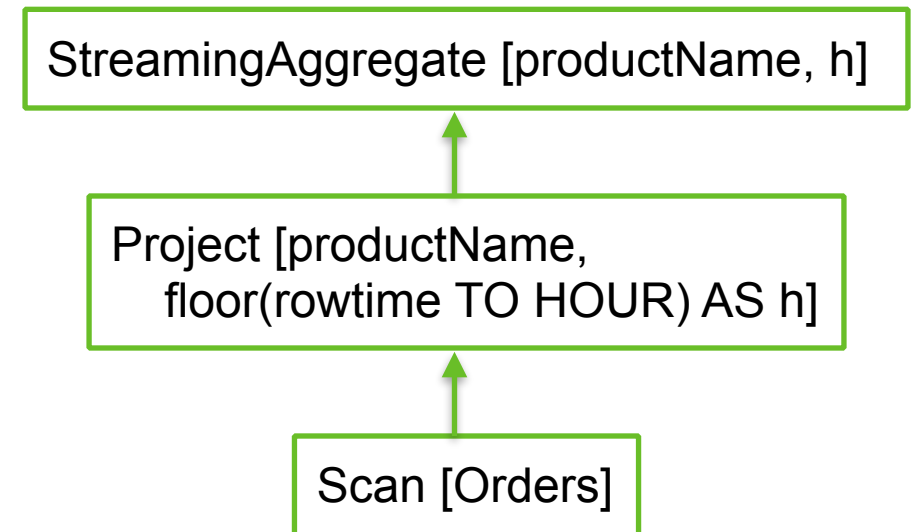


# Streaming - efficient implementation

```
SELECT STREAM DISTINCT productName,  
       floor(rowtime TO HOUR) AS h  
FROM Orders
```

**Can create efficient implementation:**

- **Input is sorted by timestamp**
- **Only need to aggregate an hour at a time**
- **Output timestamp tracks input timestamp**
- **Therefore it is safe to cancel out the Chi and Delta operators**



# Algebraic transformations - streaming

$\text{delta}(\text{filter}(c, R)) \rightarrow \text{filter}(\text{delta}(c, R))$

$\text{delta}(\text{project}(e1, \dots, en, R)) \rightarrow \text{project}(\text{delta}(e1, \dots, en, R))$

$\text{delta}(\text{union}(R1, R2)) \rightarrow \text{union}(\text{delta}(R1), \text{delta}(R2))$

$$(f + g)' = f' + g'$$

$\text{delta}(\text{join}(R1, R2, c)) \rightarrow \text{union}(\text{join}(R1, \text{delta}(R2), c),$   
 $\text{join}(\text{delta}(R1), R2), c)$

$$(f \cdot g)' = f \cdot g' + f' \cdot g$$

Delta behaves like “differentiate” in differential calculus,  
Chi like “integrate”.

# Apache Calcite



# Apache Calcite

## Apache incubator project since May, 2014

- Originally named Optiq

## Query planning framework

- Relational algebra, rewrite rules, cost model
- Extensible


## Packaging

- Library (JDBC server optional)
- Open source
- Community-authored rules, adapters


## Adoption

- **Embedded:** Lingual (SQL interface to Cascading), Apache Drill, Apache Hive, Kylin OLAP, Apache Phoenix, Apache Samza
- **Adapters:** Splunk, Spark, MongoDB, JDBC, CSV, JSON, Web tables, In-memory data

# Conventional DB architecture



# Calcite architecture





# Calcite – APIs and SPIs

## Relational algebra

RelNode (operator)

- Scan
- Filter
- Project
- Union
- Aggregate
- ...

RelDataType (type)

RexNode (expression)

RelTrait (physical property)

- RelConvention (calling-convention)
- RelCollation (sort-order)
- RelDistribution (partitions)

## SQL parser

SqlNode

SqlParser

SqlValidator

## Metadata

Schema

Table

Function

- TableFunction
- TableMacro

Lattice

## JDBC driver

## Transformation rules

RelOptRule

- MergeFilterRule
- PushAggregateThroughUnionRule
- 100+ more

Global transformations

- Unification (materialized view)
- Column trimming
- De-correlation
- Join ordering

## Cost, statistics

RelOptCost

RelOptCostFactory

RelMetadataProvider

- RelMdColumnUniqueness
- RelMdDistinctRowCount
- RelMdSelectivity

# Data independence

**A core principle of data management**

**Data independence is a contract:**

- **Applications do not make assumptions about the location or organization of data**
- **The DBMS chooses the most efficient access path**

**Requires:**

- **Declarative query language**
- **Query planner**

**Allows:**

- **The DBMS (or administrator) can re-organize the data without breaking the application**
- **Redundant copies of the data (indexes, materialized views, replicas)**
- **Novel algorithms**
- **Novel data formats and organizations (e.g. b-tree, r-tree, column store)**

# Hadoop

Name  
node  
(HDFS)

Application  
master  
(YARN)

Zookeeper

CPU

Memory

Disk

B1

B2

B3


B4

# Hadoop scales

Commodity hardware

Storage, memory and CPU all scale as you add nodes

N replicas of each block (typically 3) give redundancy & scheduling flexibility




# Hadoop query execution

Data flow among operators running on nodes

Nodes are assigned to work on blocks that have a replica locally

Memory is used for file blocks and for scratch space (e.g. hash tables)



# Data independence and Hadoop

**Hadoop is very flexible when data is loaded**

**That flexibility has made it hard for the system to optimize access**

**Materialized views are an opportunity to “crack” the data, and create copies in other formats**

# Calcite: Lattices and tiles

## Materialized view

A table whose contents are guaranteed to be the same as executing a given query.

## Lattice

Recommends, builds, and recognizes summary materialized views (tiles) based on a star schema.

A query defines the tables and many:1 relationships in the star schema.

## Tile

A summary materialized view that belongs to a lattice.

A tile may or may not be materialized.

Materialization methods:

- Declare in lattice
- Generate via recommender algorithm
- Created in response to query


(FAKE SYNTAX)

```
CREATE MATERIALIZED VIEW t AS
SELECT * FROM Emps
WHERE deptno = 10;
```

```
CREATE LATTICE star AS
SELECT *
FROM Sales AS s
JOIN Products AS p ON ...
JOIN ProductClasses AS pc ON ...
JOIN Customers AS c ON ...
JOIN Time AS t ON ...;
```

```
CREATE MATERIALIZED VIEW zg IN star
SELECT gender, zipcode,
       COUNT(*), SUM(unit_sales)
FROM star
GROUP BY gender, zipcode;
```

# Tiled, in-memory materializations



Where we're going... algebraic cache: <http://hortonworks.com/blog/dmmq/>



# Summary

- 1. Relational algebra allows us to reason about queries, and is the foundation of query planning**
- 2. Hadoop is deconstructing the DBMS, and enabling new languages, engines and data formats**
- 3. Data independence is more important than ever**
- 4. Apache Calcite - an implementation of relational algebra**

# Thank you!



@julianhyde

<http://calcite.incubator.apache.org>